

Generalized Linear Models

Melissa Guzman

March 16, 2016

Etherpad:

<https://public.etherpad-mozilla.org/p/404GLMs>

Link to all material:


<https://github.com/Imguzman/GLMs>

**What is a
frequency
distribution?**


**What frequency
distributions
can you
remember?**

Frequency Distributions


Normal




Binomial



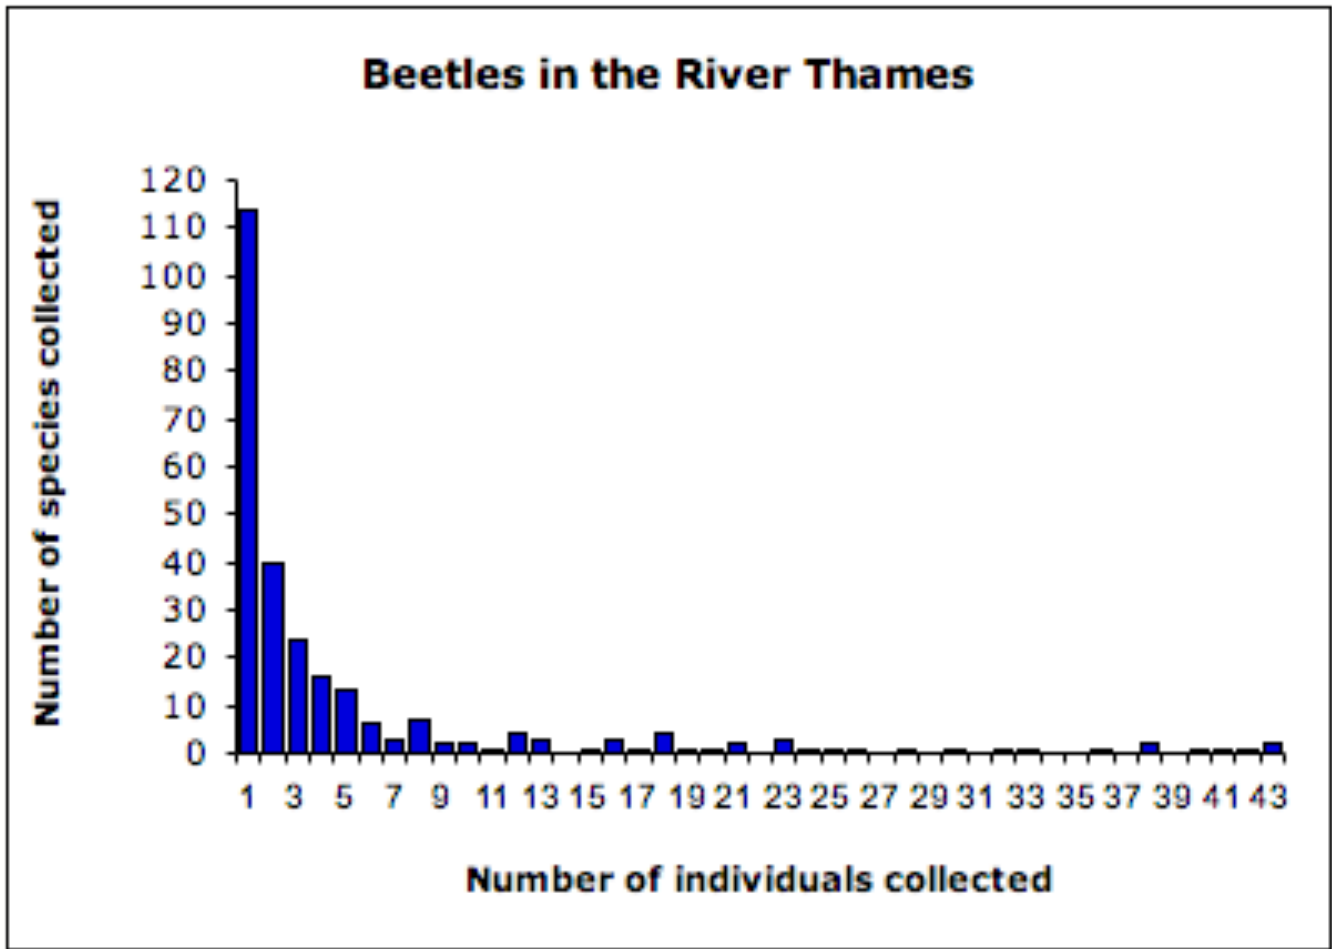
Poisson



Negative Binomial



Why do we care?



Relative species abundance of beetles sampled
from the river Thames collected by C.B.
Williams (1964). (Magurran 2004)

Learning Objectives

By the end of this lesson the students will:

- Differentiate and categorize GLMs vs ANOVAS and regressions
- Identify the different components of GLMs
- Select the most appropriate GLM
- Carry out a GLM

Generalised Linear Models (GLM)

Linear Models (LM)

“Errors” are normally distributed

Multiple regression

ANOVA

t-tests

“Errors” are not normally distributed

e.g.

- Binomial
- Negative binomial
- Poisson
- etc.

LO: Differentiate and categorize GLMs vs ANOVAS and regressions

When do we use them?

Used when the *residuals* from a linear model are not-normal

Whilst it is the error distribution (the residuals) that is important, the distribution of the dependent variable has a strong influence on this.

General formula

$$y = g(\beta_0 + \beta_1 X_1 + \beta_2 X_2 + \dots + \beta_i X_i) + \varepsilon$$

X are your explanatory variables

y your response variable

β are the coefficients

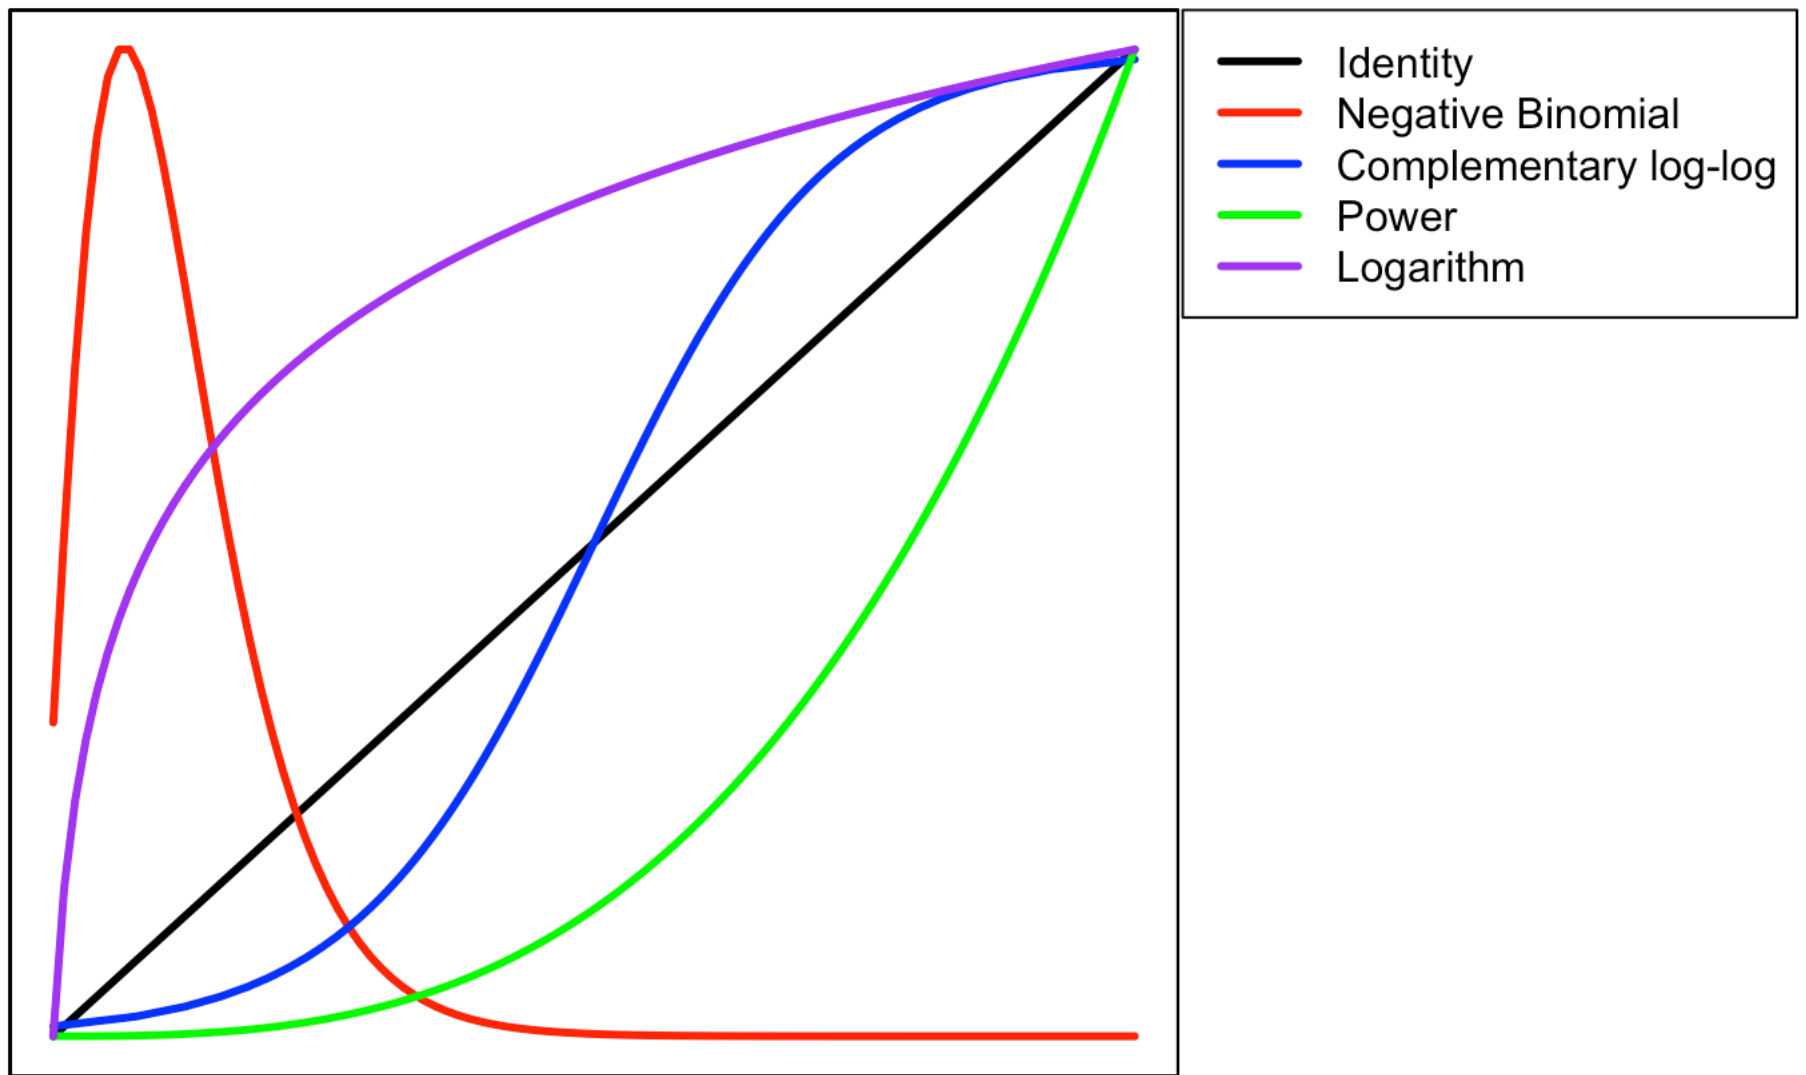
ε is the error term

Where g is a function (called the link function) which transforms each value of y in relation to the linear predictors (the variables and their coefficients) i.e. the link function transforms the dependent variable within the model.

LO: Identify the different components of GLMs


What is the link function?

Link Functions: defining the shape of the relationship between the dependent & independent variables.




Why use a link function instead of transforming the data?

Original data




Log transformed




Let's fit some models to that data

```
model1 <- lm(log(y)~x)
model2 <- glm(y~x, family = gaussian (link = "log"), start=c(1, 1))
```

LM



GLM



Using a link function is often better than simply transforming the dependent variable, because it doesn't simply attempt to make the variance constant. Instead, it fits a function that “adjusts” the variance as the data is being fitted.

What links should I use?

Families

gaussian (c)

inverse.gaussian (c)

gamma (c)

quasi (c)

poisson (d)

quasipoisson (d)

negative binomial (d)

binomial (d)

quasibinomial (d)

d = discrete, c = continuous

First link is the default

Links

identity, log, inverse

$1/\mu^2$, inverse, identity, log

inverse, identity, log

logit, probit, cloglog, identity, inverse,
 $1/\mu^2$, sqrt, power

log, identity, sqrt

identity, logit, probit cloglog

log, sqrt, identity *use glm.nb function from MASS library*

logit, probit, cauchit, log, cloglog

identity, logit, probit, cloglog

So many options... How do I choose?

Start with the default link function for the family error term that you have chosen. E.g. for family = binomial, the “logit” is the default link function.


The “correct” combination of family and link function is arrived at by trial and error, comparing

- residual distributions (aiming for normal)
- adjusted R² values (where relevant, higher is better)
- AIC values (lower is better) of different nested models


LO: Select the most appropriate GLM

Challenge


How old are you?




How many siblings do you have?




How tall are you?



What is your sex?



How many nights a week do you sleep at least 6 hours



What type of distribution are these?

Can you think of other types of data with the same distribution?

What family and link functions would you use?

How do I do it in R?

Let’s look at some data first.


```
Lizards <- read.csv('Data/lizards.txt', sep = "\t")
head(Lizards)
```

##	n	sun	height	perch	time	species
## 1	20	Shade	High	Broad	Morning	opalinus
## 2	13	Shade	Low	Broad	Morning	opalinus
## 3	8	Shade	High	Narrow	Morning	opalinus
## 4	6	Shade	Low	Narrow	Morning	opalinus
## 5	34	Sun	High	Broad	Morning	opalinus
## 6	31	Sun	Low	Broad	Morning	opalinus

LO: Carry out a GLM

hist(Lizards\$n)

Histogram of Lizards\$n







```
LizardsModelLM <- lm(n ~ sun + height + species, data = Lizards)
```

```
LizardsModel <- glm(n ~ sun + height + species, family = poisson (link = log), data = Lizards)
```

Check assumptions of the model





I. Distribution of the residuals

```
par(mfrow = c(2,2))  
plot(LizardsModelLM)
```



Let's see the GLM

```
par(mfrow = c(2,2))  
plot(LizardsModel)
```



2. Overdispersion

In a poisson distribution, mean = variance

Overdispersion parameter $\theta = \frac{\text{residual deviance}}{\text{residual degrees of freedom}}$

```
LizardsModel$deviance / LizardsModel$df.residual
```

```
## [1] 6.352926
```

If your model is overdispersed use the *quasi* family. In this case, quasipoisson

pseudo R^2 explained deviance

Pseudo $R^2 = \frac{\text{null deviance} - \text{residual deviance}}{\text{null deviance}}$

```
(LizardsModel$null.deviance - LizardsModel$deviance) / LizardsModel$null.deviance
```

```
## [1] 0.6210064
```

Model Summary

summary(LizardsModel)

```
##
## Call:
## glm(formula = n ~ sun + height + species, family = poisson(link = log),
##      data = Lizards)
##
## Deviance Residuals:
##      Min       1Q   Median       3Q      Max
## -5.9107  -1.6984  -0.6075   0.6341   6.2438
##
## Coefficients:
##              Estimate Std. Error z value Pr(>|z|)
## (Intercept)    0.97277    0.12778   7.613 2.68e-14 ***
## sunSun         1.48684    0.10858  13.694 < 2e-16 ***
## heightLow     -0.60659    0.08812  -6.884 5.83e-12 ***
## speciesopalinus 1.17576    0.09919  11.853 < 2e-16 ***
## ---
## Signif. codes:  0 '***' 0.001 '**' 0.01 '*' 0.05 '.' 0.1 ' ' 1
##
## (Dispersion parameter for poisson family taken to be 1)
##
##      Null deviance: 737.56  on 47  degrees of freedom
## Residual deviance: 279.53  on 44  degrees of freedom
## AIC: 450.78
##
## Number of Fisher Scoring iterations: 5
```

Challenge

Using the Boar csv, run the appropriate GLM.

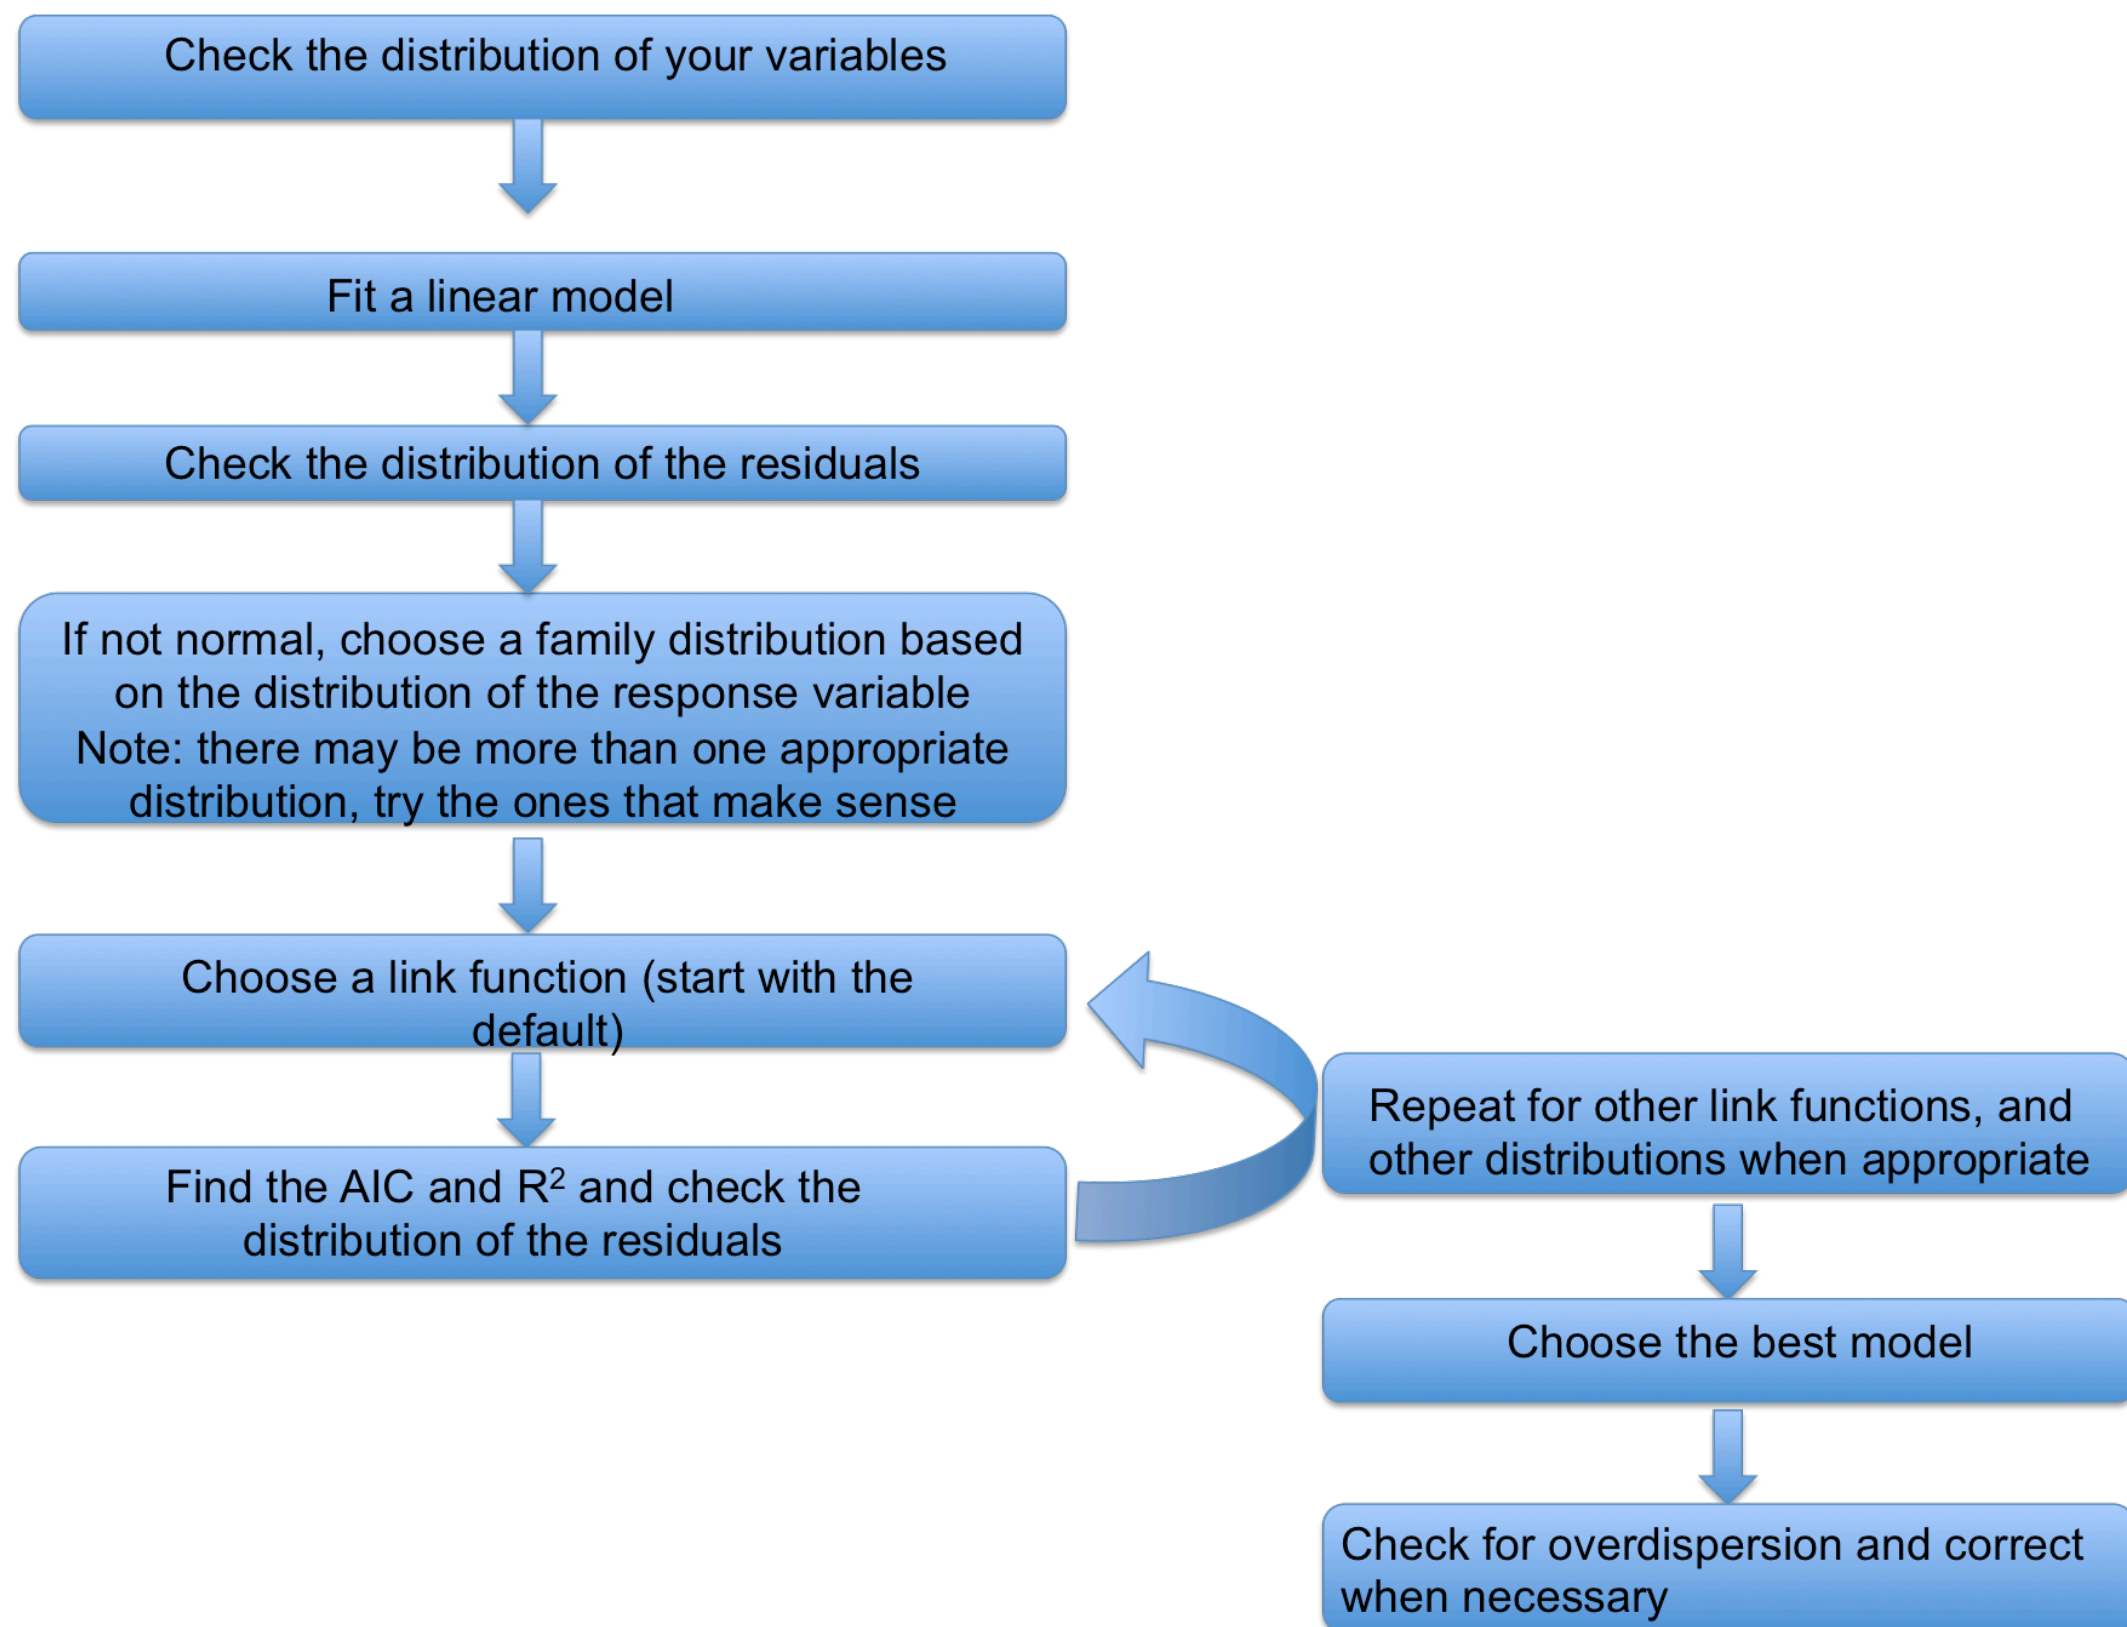
Make sure you check the model assumptions and all appropriate link functions

Learning Objectives

By the end of this lesson the students will:

- Differentiate and categorize GLMs vs ANOVAS and regressions
- Identify the different components of GLMs
- Select the most appropriate GLM
- Carry out a GLM

Summary



Challenge 2 Solution

Look at the data first

```
Boar <- read.csv('Data/Boar.csv')

head(Boar)
```

##	Tb	sex	age	length
## 1	0	1	1	46.5
## 2	0	2	1	47.0
## 3	0	1	1	48.0
## 4	0	1	1	51.5
## 5	0	2	1	53.0
## 6	0	2	1	53.0

Tb is the response variable. Sex, age and length are explanatory variables.


```
summary(Boar)
```

##	Tb	sex	age	length
## Min.	:0.0000	Min. :1.000	Min. :1.000	Min. : 46.5
## 1st Qu.	:0.0000	1st Qu.:1.000	1st Qu.:3.000	1st Qu.:106.0
## Median	:0.0000	Median :2.000	Median :3.000	Median :121.0
## Mean	:0.4758	Mean :1.582	Mean :3.142	Mean :116.8
## 3rd Qu.	:1.0000	3rd Qu.:2.000	3rd Qu.:4.000	3rd Qu.:129.5
## Max.	:1.0000	Max. :2.000	Max. :4.000	Max. :165.0
## NA's	:58	NA's :35	NA's :7	NA's :110

Let’s check the distribion to see how it looks like.

```
hist(Boar$Tb)
```





Histogram of Boar\$Tb



Looking at the summary and histogram, Tb is a categorical variables with two categories. Therefore we need to use the binomial distribution.

We can start with a linear model

```
LinearBoarModel <- lm(Tb ~ sex * age * length, data = Boar)
par(mfrow = c(2,2))
plot(LinearBoarModel)
```



The residuals don't look very normally distributed. Let's use a binomial family.

The default link functions for the binomial distribution are: logit, probit, cauchit, log, cloglog Let's start with the logit function

```
LogitBoarModel <- glm(Tb ~ sex * age * length, family = binomial (link = logit), data = Boar)
```

Let's look at the AIC and pseudo R^2

```
AIC(LogitBoarModel)
```





```
## [1] 653.3006
```

```
(LogitBoarModel$null.deviance - LogitBoarModel$deviance) / LogitBoarModel$null.deviance
```

```
## [1] 0.06451869
```

And the distribution of the residuals

```
par(mfrow = c(2,2))
plot(LogitBoarModel)
```



Look at all the other link functions and choose the one with the lowest AIC and better distribution of the residuals.

```
probitBoarModel <- glm(Tb ~ sex * age * length, family = binomial (link = probit), data = Boar)

cauchitBoarModel <- glm(Tb ~ sex * age * length, family = binomial (link = cauchit), data = Boar)

logBoarModel <- glm(Tb ~ sex * age * length, family = binomial (link = log), data = Boar)

cloglogBoarModel <- glm(Tb ~ sex * age * length, family = binomial (link = cloglog), data = Boar)

AIC(LogitBoarModel, probitBoarModel, cauchitBoarModel, logBoarModel, cloglogBoarModel)
```

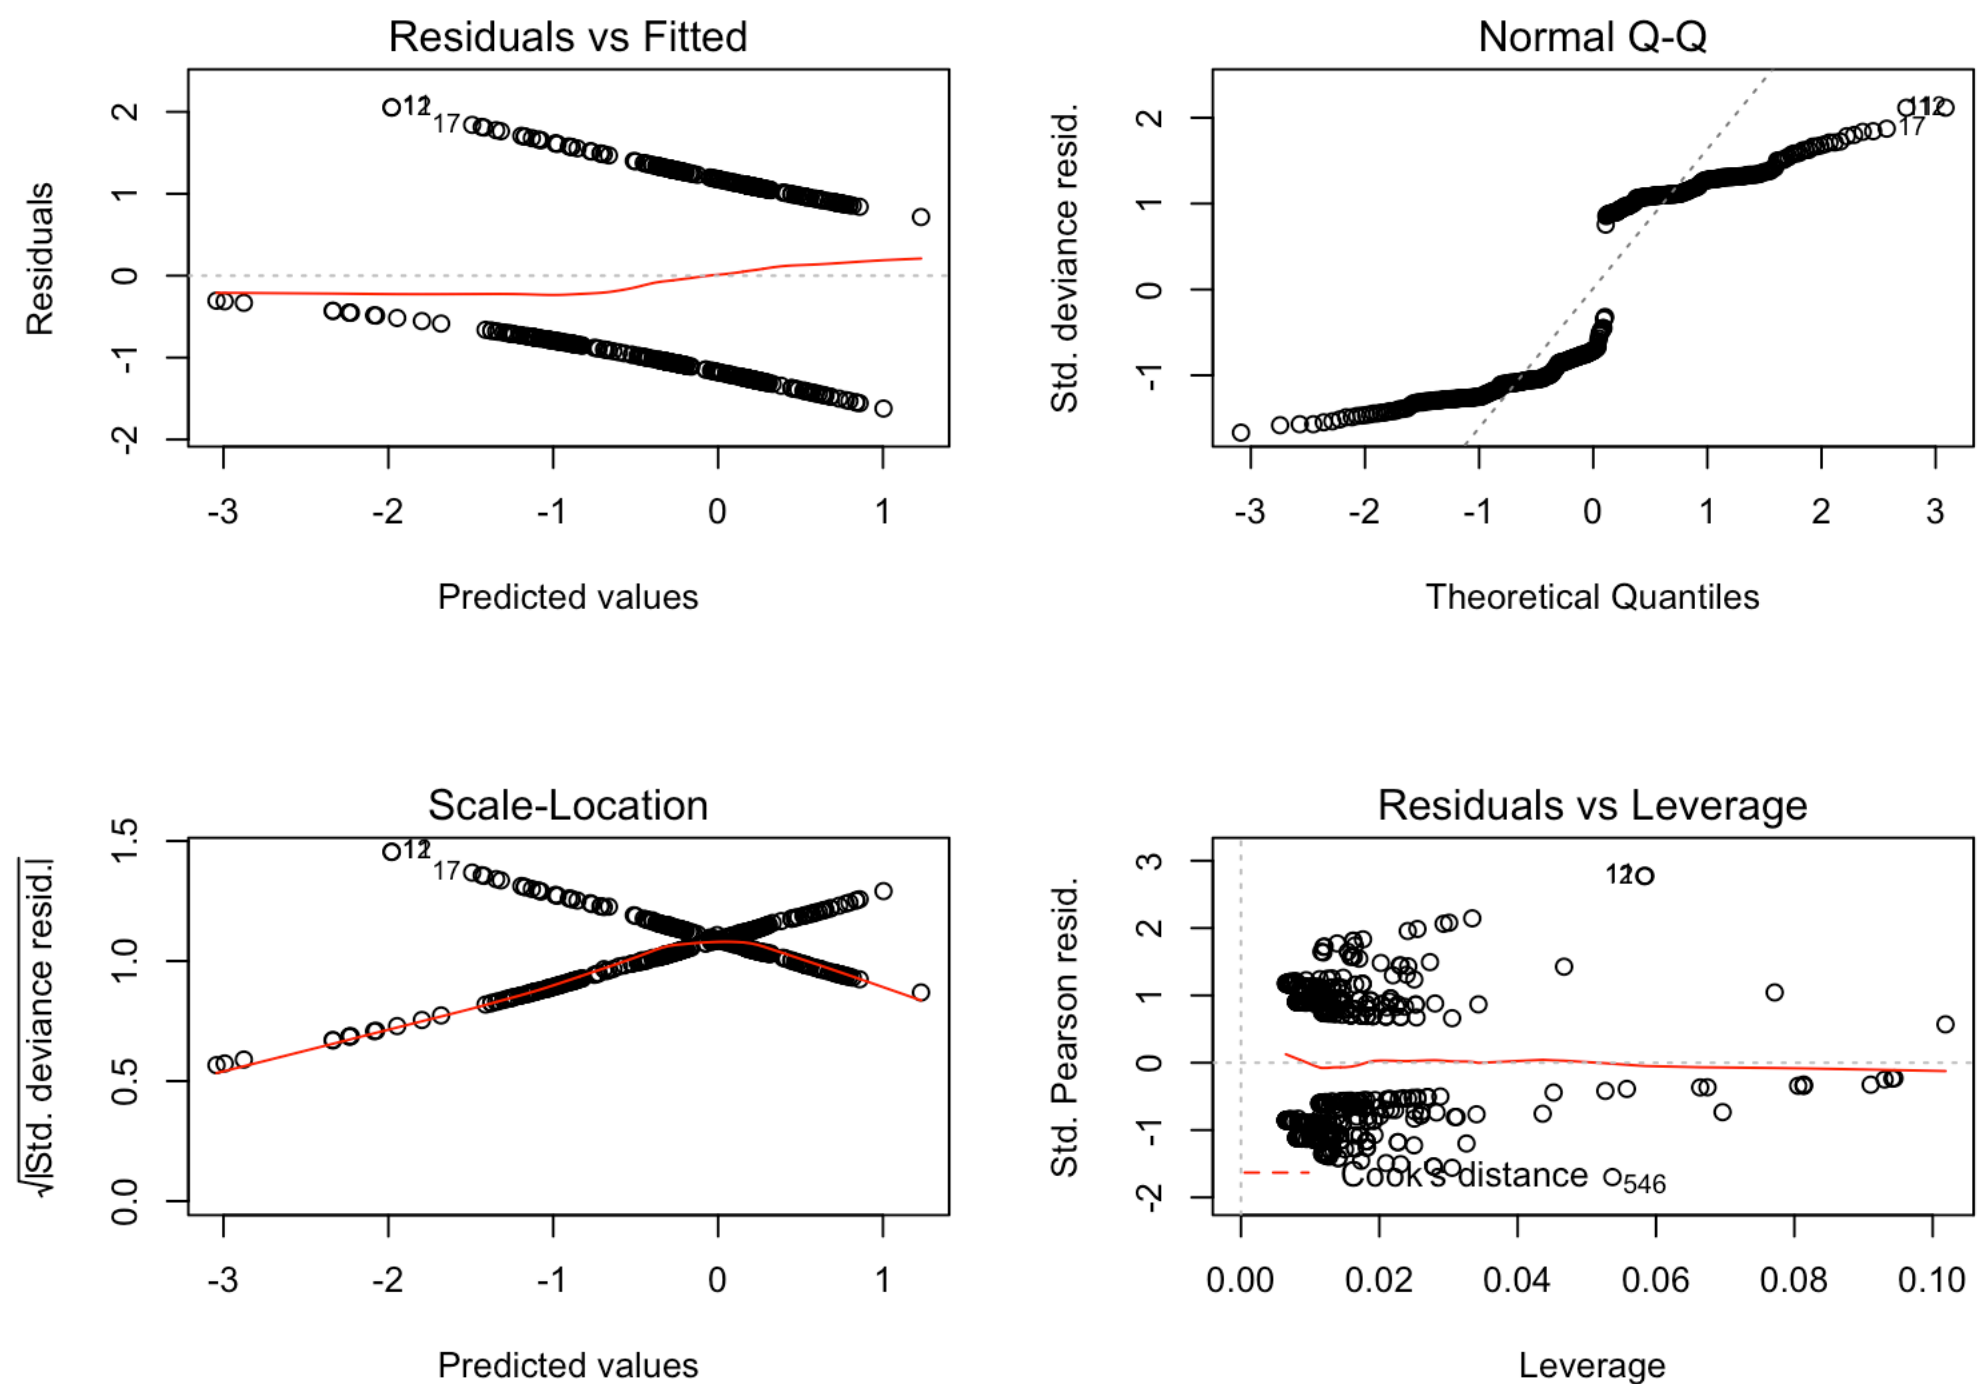
```
##           df      AIC
## LogitBoarModel    8 653.3006
## probitBoarModel   8 653.3818
```

```
## cauchitBoarModel 8 652.2948
## logBoarModel      8 652.8498
## cloglogBoarModel  8 653.1425
```

Logit has the lowest AIC, so let's keep working with that

Let's check the distribution of the residuals.

```
par(mfrow = c(2,2))
plot(LogitBoarModel)
```



We also need to check for overdispersion

```
LogitBoarModel$deviance / LogitBoarModel$df.residual
```

```
## [1] 1.311318
```

θ is close to 1, so we can keep the binomial family.

We can keep refining the model using model simplification to choose a model we are happy with.

